

3--- title: An Automation Development Stack theme:  
night revealOptions: transition: 'fade'

# AN AUTOMATION DEVELOPMENT STACK

One implementation of an automation and configuration stack that allows users to leverage a complex stack via a **high level of abstraction**.

Christopher Steel, Systems Developer (MCIN)  
December 14th, 2018

# OBTAINING A HIGH LEVEL OF ABSTRACTION

In this case, obtaining a high level of abstraction:

- Makes it "easy" to leverage the magical powers of some complex technology
- Allows for the creation of provider agnostic automation scripts which in turn
- Leads to standardized automation roles
- Likely to increase development velocity, perhaps significantly

# NOT EXACTLY MAGIC BUT...

"Any sufficiently advanced technology is indistinguishable from magic."

Arthur C. Clarke's Third Law

- No rabbits, just a software stack and
- A few virtual machines and **containers** and
- Perhaps a **COW** or two?!

# THE DEVELOPMENT STACK

- Ansible - automation software
- LXC / LXD - containers and hypervisor
- Molecule - Ansible role development aid
- OpenZFS - powerful file system and logical volume manager
- Vagrant -
- VirtualBox - open source virtual machine application

# ANSIBLE

An open source automation engine

- Easy to **grok** SERIOUSLY???
- Provisioning and configuration management
- Application deployment and intra-service orchestration
- Modules in abundance made using any language
- SSH / No agents required

# LXC AND LXD

## Linux **System** Containers

- As fast as bare metal (2% slower)
- Ultralight hypervisor
- Well considered CLI / REST API
- **Magical** properties

# MOLECULE

- Allows for consistently developed Ansible roles
- Supports **multiple**:
  - instances
  - operating systems
  - distributions
  - test frameworks
  - testing scenarios
  - virtualization providers / target nodes (bare-metal, VMs, cloud(s) and/or containers)



# OPENZFS

OpenZFS acts as both a volume manager and a file system it has a few magical properties as well.

- Copy-on-write (COW) transactional object model
- Continuous integrity checking and automatic repair
- Lots of other interesting stuff, RAID-Z, Native NFSv4 ACLs...
- Ubuntu has a really nice OpenZFS, FreeNAS an appliance version.

# VIRTUALBOX

- Open source
- Well understood
- Mature
- Runs on many OS, able to host many OS...

# VIRTUALIZATION OVERVIEW

Before we get started with the details of this particular development stack, lets take some 3time and talk a bit about virtualization and how to go about choosing the right type of virtualization for your particular project or experiment.

# REASONS YOU MIGHT WANT (OR NEED) TO VIRTUALIZE

- Isolation / security
- Density / effective use of resources
- Efficiency and/or speed
- **Repeatability** and/or portability
- Required in order to...
- Other reasons

# HYPERVERSORS

What is a Hypervisor? Generally speaking a  
Hypervisor:

- Allows a physical **host** to operate multiple VMs or containers as **guests**
- Is **a process** that separates a system's OS and apps from the physical hardware
- Is **software**, but could be **embedded** software on a chip

# MAJOR TYPES OF VIRTUALIZATION (HYPERVISORS)

- Type 1 - **Native** ( bare metal )
- Type 2 - **Hosted** ("on top of" OS)
- Type C - **OS level virtualization / Containerization / Other**

# TYPE 1 - **NATIVE** ("BARE METAL")

Hardware <--> Hypervisor <--> Hosted OS

- Runs directly on the host's hardware
- Controls the hardware
- Manages the guest operating system(s)

## TYPE 1 (BARE METAL) HYPERVISOR EXAMPLES:

- VMware ESXi
- Xen ( open source - loads it's own paravirtualized host operating system)
- Xbox One system software



## SOME SPECIAL PROPERTIES OF TYPE 1 VIRTUALIZATION (HYPERVISORS)

- Speed and Flexibility?
  - Windows hosts (Xen / VMware)
  - FreeBSD (Xen)
  - NetBSD (Xen)
- Security implications
  - No host OS(!)
  - Probably a full OS on your VM's

# TYPE 2 - **HOSTED** HYPERVISORS

Hardware <--> Host OS <--> Hypervisor <--> Hosted OS

- Generally an application that runs on a conventional (host) operating system OS
- The guest operating system runs as a process on the host OS
- Abstracts the guest OS from the host OS

## TYPE 2 (HOSTED) HYPERVISOR EXAMPLES:

- KVM
- VirtualBox
- VMware Workstation Pro

## SOME PROPERTIES OF TYPE 2 VIRTUALIZATION

- Isolation of instances
- Support for non-linux OS
- Well understood, familiar to more users

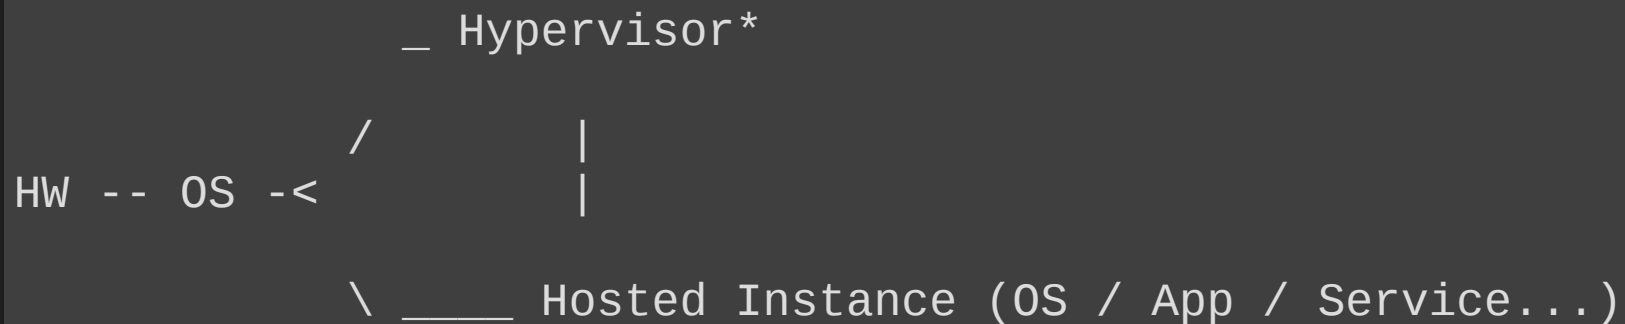
# TYPE C VIRTUALIZATION

Catch all including: containers, linux system containers, jails and chroot jails, AIX Workload Partitions (WPARs) and virtual environments.

- Diverse but shared history
- May make use of a hypervisor
- May allow for the existence of multiple isolated user-space instances
- May be recursively virtualizable

# CONTAINERS, TYPE C ARCHITECTURES THAT EMPLOY A HYPERVISOR

Instances (launched containers) look like real computers from the point of view of the programs running in them.:



## EXAMPLES OF TYPE C VIRTUALIZATION USING HYPERVISORS:

- Docker - Used to be an alternative Hypervisor for using Linux System Containers (LXC)
- LXC (Linux System Containers) - LXD is an extension of the LXC hypervisor

# CONTAINERS

Containers have some really nice features.

- Usually faster and lighter than virtual machines.
- Linux support



# SELECTING THE RIGHT CONTAINER FOR THE JOB

- **Usage** (microservice, enterprise application, development, devops, HPC?)
- **Users** (are you a developer, admin, power user or a researcher?)
- **Support** & Community, Popularity
- **Strengths & Weaknesses**
- **Sharing containers** - Image/Container Registry

# SINGULARITY

Designed for HPC

- Can be "run" or executed directly by file name
- Image-based containers
- Must be root outside of container to become root inside of container
- No root daemon owned processes
- Can import docker containers
- Learning curve

## SINGULARITY USERS

- MCIN users
- Calcul Quebec / Calcul Canada
- HPC all over the world

# SINGULARITY PIPELINE EXAMPLE

After (a lot of) careful preparation...:

```
singularity run --app cat catdog.simg
```

Output:

```
Meow , this is Cat
```

# APPLICATION CONTAINERS - DOCKER

- Designed for **isolating applications**
- Docker to **HPC** via Singularity
- Use of layers and disabling of persistence results in **lower disk IO**
- **Issues** using apps that expect cron, ssh, daemons, logging and other system stuff
- Can use copy-on-write (CoW) for images and containers

## DOCKER USERS

- MCIN
- Development and test organizations
- Many, many others

# LXC & LXD (LINUX SYSTEM CONTAINERS)

- Acts like a **normal OS environment**: hostname, IP address, file systems, init.d, SSH access
- Nearly **as fast as bare metal**, parallelism possible
- Efficiently run **one or more multi-process applications**
- Linux-native, can leverage **CoW** using ZFS backing.
- Windows (10 with Linux subsystem enabled) and OSX clients\*

## WHO USES LXC / LXD

- IT Operators / DevOps
- MCIN
- Most Canonical Websites



## LXC/LXD USAGE

- Will cover this in the next section

# CHOOSING YOUR VIRTUALIZATION TARGET(S)

- Do you have / require root access on your virtualization project?
- What will you be virtualizing? An application, many applications, part of a pipeline?
- Host - Where will your container be running?
- Which OSs will be running on your host and guest systems?

# MOLECULE

Enables the development of **provider agnostic** Ansible playbooks and roles.

- Includes built-in **provider specific Ansible tasks** created using Ansible provider modules
- Allows for the creation of **custom providers** when Ansible provider modules are not available

# MOLECULE INIT

- **molecule init role** is used to create a new **provider agnostic Ansible role** and a default scenario
- **molecule init scenario** is used to create additional scenarios for an existing Ansible role

```
molecule init role -r myrole  
cd myrole  
rm -R molecule/default  
molecule init scenario -s default -d vagrant -r myrole  
molecule init scenario -s lxd -d lxd -r myrole
```

# GENERATED FILES

Generating a **scenario** creates a scenario directory and contents. The `INSTALL.rst` file includes information on any additional requirements molecule has for supporting a particular scenario:

```
ls -al molecule/default/  
-rw-r--r-- 1 cjs cjs 260 Dec 5 20:08 INSTALL.rst  
-rw-r--r-- 1 cjs cjs 303 Dec 5 20:23 molecule.yml  
-rw-r--r-- 1 cjs cjs 64 Dec 5 20:08 playbook.yml  
-rw-r--r-- 1 cjs cjs 235 Dec 5 20:08 prepare.yml  
drwxr-xr-x 2 cjs cjs 4096 Dec 5 20:08 tests
```

# MOLECULE.YML

- Each scenario has a `molecule.yml` file
- By default, a single instance is created called **instance**
- You may want to call it something more meaningful

```
nano molecule/default/molecule.yml  
nano molecule/1xd/molecule.yml
```

# MOLECULE CREATE

I use vagrant/virtualbox and lxd/lxd-based scenarios for most of my testing. I configured LXC/LXD to be backed with ZFS (as a file). So once I have a local copy of the containers base image when ever **molecule create** is run against an lxd scenario it uses **Copy On Write** (COW) to create the instance as I configured LXC/LXD to be backed using ZFS as a file. (LAST SENTENCE NEEDS SERIOUS AUTHOR REVIEW)

```
# molecule check -s lxd      # do a dry run
time molecule create          # create default instance
time molecule create -s lxd   # create lxd instance (using COW)
```

# OTHER MOLECULE COMMANDS

Molecule has a total of 16 high-level commands at this time. Here are three:

```
time molecule converge -s lxd # configure the instance(s) using  
time molecule test -s lxd     # Run all tests...
```



# MOLECULE LIST

molecule list

# MOLECULE DESTROY

molecule destroy

# LXC/LXD

Our storage pool

```
lxc storage info lxd
```

# LAUNCHING A TEST CONTAINER

Checking our space usage:

```
lxc storage info lxd
```

## LIST OUR CONTAINERS

```
lxc list
```

## SOME KEY POINTS

- We have learned a lot and we are still learning about our new stack.
- It makes a lot of things that used to be impractical practical.
- **High levels of abstraction** rock.
- Use the right kind of virtualization for the job at hand.
- For extensive testing, you will want to leverage the fastest appropriate provider driver. (NEEDS AUTHOR REVIEW)
- Use one or more scenarios per driver depending on your needs.

# IN CLOSING

Raising the level of abstraction makes automation much easier as it hides many levels of complexity. In this stack, this includes:

- LXC and LXD enabling the configuration and leveraging of a ZFS back end.
- molecule allowing use to ignore provider specifics once the provider and instance is configured.  
(NEEDS AUTHOR REVIEW)

This in turn:

- Allows us to provisioning to multiple targets. (NEEDS AUTHOR REVIEW)
- Allows us to develop target agnostic Ansible roles.
- Enables a significant increase in development velocity.

**REFERENCES**

**COMPARISONS**

**WEB PAGES**

- <https://robin.io/blog/linux-containers-comparison-lxc-docker/>
- <https://robin.io/blog/containers-deep-dive-lxc-vs-docker-comparison/>
- <https://www.xenproject.org/users/virtualization.html>

## **PAPERS**

- Formal Requirements for Virtualizable Third Generation Architectures-10.1.1.141.4815
- Analysis of Virtualization Technologies for High Performance Computing Environments
- Performance Evaluation of Container-based Virtualization for High Performance Computing Environments
- NIST.SP.800-125A-F - Security Recommendations for Hypervisor Deployment
- VIRTUALIZATION TECHNIQUES & TECHNOLOGIES: STATE-OF-THE-ART

**PRESENTATIONS**

**MCIN**